

Analyze_AB_test_results_notebook

January 22, 2019

0.1 Analyze A/B Test Results

0.2 Table of Contents

- Introduction
- Part I - Probability
- Part II - A/B Test
- Part III - Regression

0.2.1 Introduction

In this project, we will be working to understand the results of an A/B test run by an e-commerce website. The goal is to help the company understand if they should implement the new page, keep the old page, or perhaps run the experiment longer to make their decision.

Part I - Probability To get started, We'll import necessary libraries.

```
In [169]: import pandas as pd
import numpy as np
import random
import matplotlib.pyplot as plt
%matplotlib inline
#We are setting the seed to get the same sequence of numbers
random.seed(42)
```

- Now, we'll read in the `ab_data.csv` data. Store it in `df`

```
In [170]: # read in the dataset and take a look at the top few rows
df = pd.read_csv('ab_data.csv')
df.head()
```

```
Out[170]:
```

	user_id	timestamp	group	landing_page	converted
0	851104	2017-01-21 22:11:48.556739	control	old_page	0
1	804228	2017-01-12 08:01:45.159739	control	old_page	0
2	661590	2017-01-11 16:55:06.154213	treatment	new_page	0
3	853541	2017-01-08 18:28:03.143765	treatment	new_page	0
4	864975	2017-01-21 01:52:26.210827	control	old_page	1

```
In [171]: df.timestamp.max(), df.timestamp.min()
```

```
Out[171]: ('2017-01-24 13:41:54.460509', '2017-01-02 13:42:05.378582')
```

```
In [172]: # find the number of rows in the dataset
df.shape
```

```
Out[172]: (294478, 5)
```

```
In [173]: # The number of unique users in the dataset
df.user_id.nunique()
```

```
Out[173]: 290584
```

```
In [174]: # The proportion of users converted
df['converted'].mean()
```

```
Out[174]: 0.11965919355605512
```

```
In [175]: # The number of times the `new_page` and `treatment` don't line up
# Either new_page with control OR old_page with treatment
(df.query('landing_page == "new_page" and group == "control").shape[0] +
df.query('landing_page == "old_page" and group == "treatment").shape[0])
```

```
Out[175]: 3893
```

```
In [176]: # Check if there are any missing values
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 294478 entries, 0 to 294477
Data columns (total 5 columns):
user_id      294478 non-null int64
timestamp    294478 non-null object
group        294478 non-null object
landing_page  294478 non-null object
converted     294478 non-null int64
dtypes: int64(2), object(3)
memory usage: 11.2+ MB
```

- Now, we'll handle rows where treatment is not aligned with new_page or control is not aligned with old_page

```
In [177]: # Create a new dataset that meets the specifications and Store it in `df2`
df0 = df.query('landing_page == "old_page" and group != "treatment"')
df0.shape[0]
```

```
Out[177]: 145274
```

```
In [178]: df1 = df.query('landing_page == "new_page" and group != "control"')
df1.shape[0]
```

```
Out[178]: 145311
```

```
In [179]: df2 = df0.append(df1)
df2.head()
```

```
Out[179]:
```

	user_id	timestamp	group	landing_page	converted
0	851104	2017-01-21 22:11:48.556739	control	old_page	0
1	804228	2017-01-12 08:01:45.159739	control	old_page	0
4	864975	2017-01-21 01:52:26.210827	control	old_page	1
5	936923	2017-01-10 15:20:49.083499	control	old_page	0
7	719014	2017-01-17 01:48:29.539573	control	old_page	0

```
In [180]: # Double Check all of the correct rows were removed - this should be 0
(df2[((df2['group'] == 'treatment') ==
(df2['landing_page'] == 'new_page')) == False].shape[0])
```

```
Out[180]: 0
```

```
In [181]: # Unique `user_id`s are in `df2`
df2.user_id.nunique()
```

```
Out[181]: 290584
```

```
In [182]: # Find the `user_id` repeated in `df2`
df2['user_id'].value_counts()
```

```
Out[182]:
```

773192	2
630732	1
811737	1
797392	1
795345	1
801490	1
799443	1
787157	1
793302	1
817882	1
842446	1
815835	1
805596	1
803549	1
809694	1
807647	1
895712	1
840399	1
836301	1
899810	1
834242	1
936604	1
934557	1

```

940702    1
938655    1
830144    1
828097    1
832195    1
838348    1
821956    1
..
734668    1
736717    1
730574    1
775632    1
771538    1
642451    1
773587    1
783828    1
785877    1
779734    1
781783    1
759256    1
726472    1
748999    1
746950    1
753093    1
751044    1
740803    1
738754    1
744897    1
742848    1
634271    1
632222    1
636316    1
630169    1
650647    1
648598    1
654741    1
652692    1
630836    1

```

Name: user_id, Length: 290584, dtype: int64

```

In [183]: # find the row information for the repeat 'user_id'
df2[df2['user_id'] == 773192]

```

```

Out[183]:
   user_id  timestamp      group landing_page  converted
1899  773192  2017-01-09 05:37:58.781806  treatment    new_page         0
2893  773192  2017-01-14 02:55:59.590927  treatment    new_page         0

```

```

In [184]: # Remove one of the rows with a duplicate 'user_id',
# but keeping the dataframe as 'df2'

```

```

df2.drop_duplicates(subset=['user_id'], keep='first', inplace=True)
df2[df2['user_id'] == 773192]

Out[184]:
   user_id  timestamp      group landing_page  converted
1899   773192  2017-01-09 05:37:58.781806  treatment    new_page         0

In [185]: # Find the probability of an individual converting
          # regardless of the page they receive
          (df2['converted'] == 1).mean()

Out[185]: 0.11959708724499628

In [186]: # Given that an individual was in the `control` group,
          # what is the probability they converted?
          control_prob = (df2.query('group == "control"')['converted'] == 1).mean()
          control_prob

Out[186]: 0.1203863045004612

In [187]: # Given that an individual was in the `treatment` group,
          # what is the probability they converted?
          treat_prob = (df2.query('group == "treatment"')['converted'] == 1).mean()
          treat_prob

Out[187]: 0.11880806551510564

In [188]: # Find the probability that an individual received the new page
          new = df2[df2['landing_page'] == "new_page"].count().landing_page
          new_prob = new / df2.shape[0]
          new_prob

Out[188]: 0.50006194422266881

```

The results are almost the same and so, it is hard to conclude that one page leads to more conversions or not.

0.2.2 Part II - A/B Test

Notice that because of the time stamp associated with each event, we could technically run a hypothesis test continuously as each observation was observed.

However, The hard question is do we stop as soon as one page is considered significantly better than another or does it need to happen consistently for a certain amount of time? How long do we run to render a decision that neither page is better than another?

For now, we will assume that the old page is better unless the new page proves to be definitely better at a Type I error rate of 5%. The null and alternative hypotheses will be:

```

''' Null: = ALt: >
'''

```

```

In [189]: # Find the `convert rate` for `p_{new}` under the null
          p_new = (df2['converted'] == 1).mean()
          p_new

```

Out[189]: 0.11959708724499628

```
In [190]: # Find the `convert rate` for `p_{old}` under the null
p_old = (df2['converted'] == 1).mean()
p_old
```

Out[190]: 0.11959708724499628

```
In [191]: # Find `n_{new}`
n_new = df2[df2['landing_page'] == "new_page"].count().landing_page
n_new
```

Out[191]: 145310

```
In [192]: # Find `n_{old}`
n_old = df2[df2['landing_page'] == "old_page"].count().landing_page
n_old
```

Out[192]: 145274

```
In [193]: # Simulate `n_{new}` transactions with a convert rate of `p_{new}` under the null
# Store these `n_{new}` 1's and 0's in `new_page_converted`
new_page_converted = np.random.choice(2, n_new, p_new)
```

```
In [194]: # Simulate `n_{old}` transactions with a convert rate of `p_{old}` under the null
# Store these `n_{old}` 1's and 0's in `old_page_converted`
old_page_converted = np.random.choice(2, n_old, p_old)
```


```
In [195]: # Find `p_{new}` - `p_{old}` for the simulated values
diff = new_page_converted.mean() - old_page_converted.mean()
diff
```

Out[195]: -0.0029386509123299209

```
In [196]: # Simulate 10,000 `p_{new}` - `p_{old}` values
# to the previous parts above
# Store all 10,000 values in `p_diffs`
p_diffs = (np.random.binomial(n_new, p_new, 10000)/n_new -
            np.random.binomial(n_old, p_old, 10000)/n_old)
p_diffs
```

Out[196]: array([0.00075495, 0.00138123, -0.00130314, ..., 0.00051415,
 0.00219368, -0.00073151])

```
In [197]: # Plot a histogram of the `p_diffs`
plt.hist(p_diffs);
```



```
In [198]: # Find proportion of the `p_diffs` are greater than
# the actual difference observed in `ab_data.csv`
new = df2[df2['landing_page'] == "new_page"].count().landing_page
new_prob = new / df.shape[0]
new_prob
```

```
Out[198]: 0.49344942576355449
```

```
In [199]: old = df2[df2['landing_page'] == "old_page"].count().landing_page
old_prob = old / df.shape[0]
old_prob
```

```
Out[199]: 0.49332717554452288
```

```
In [200]: actual_diff = new_prob - old_prob
actual_diff
```

```
Out[200]: 0.00012225021903161659
```

```
In [201]: df_diffs = pd.DataFrame(p_diffs)
df_diffs.columns = ['p_diffs']
```

```
In [202]: proportion=len(df_diffs.query('p_diffs > @ actual_diff'))/len(df_diffs)
proportion
```

```
Out[202]: 0.463
```

- We have calculated the number of values (differences) in our simulation greater than the actual difference (probability), this value is the P_Value and it tells us that we fail to reject the null hypothesis and that the new page is not better or the same as the old one

```
In [203]: # calculate the number of conversions for each page,
# as well as the number of individuals who received each page (using a built-in)
import statsmodels.api as sm

convert_old = df2.query('landing_page == "old_page" and converted == 1').shape[0]
convert_new = df2.query('landing_page == "new_page" and converted == 1').shape[0]
print(convert_old, convert_new)
n_old = df2[df2['landing_page'] == "old_page"].count().landing_page
n_new = df2[df2['landing_page'] == "new_page"].count().landing_page
print(n_old, n_new)
```

```
17489 17264
145274 145310
```

```
In [204]: # use `stats.proportions_ztest` to compute test statistic and p-value
z_score, p_value = (sm.stats.proportions_ztest([convert_old, convert_new],
                                                [n_old, n_new]))

z_score, p_value
```

```
Out[204]: (1.3109241984234394, 0.18988337448195103)
```

- The z-score and p-value both are greater than 0.05 which mean that we fail to reject the null hypothesis (Z score is between -1.96 and +1.96 and so the p-value is larger than 0.05, we cannot reject the null hypothesis), which means that conversion rates are not high and there is no difference between the old page and the new one. They agree with findings in parts j and k

0.2.3 Part III - A regression approach

Since the values are categorical variables we need to use the logistic regression

```
In [205]: # Create a column for the intercept
df2['intercept'] = 1
# Create `ab_page` column, which is 1 when an individual
# receives the `treatment` and 0 if `control`
df2['ab_page'] = (df2['group'] == "treatment").astype(int)
df2.head()
```

```
Out[205]:
```

	user_id	timestamp	group	landing_page	converted	\
0	851104	2017-01-21 22:11:48.556739	control	old_page	0	
1	804228	2017-01-12 08:01:45.159739	control	old_page	0	
4	864975	2017-01-21 01:52:26.210827	control	old_page	1	
5	936923	2017-01-10 15:20:49.083499	control	old_page	0	
7	719014	2017-01-17 01:48:29.539573	control	old_page	0	

	intercept	ab_page
0	1	0
1	1	0
4	1	0
5	1	0
7	1	0

```
In [206]: # Import the regression model using statsmodels
import statsmodels.api as sm
# Instantiate the model
model = sm.Logit(df2['converted'], df2[['intercept', 'ab_page']])
# fit the model
result = model.fit()
result.summary()
```

Optimization terminated successfully.
Current function value: 0.366118
Iterations 6

```
Out[206]: <class 'statsmodels.iolib.summary.Summary'>
"""
                                Logit Regression Results
=====
Dep. Variable:                converted    No. Observations:                290584
Model:                        Logit       Df Residuals:                    290582
Method:                        MLE        Df Model:                        1
Date:                         Tue, 22 Jan 2019    Pseudo R-squ.:                8.077e-06
Time:                         15:22:27         Log-Likelihood:                -1.0639e+05
converged:                     True          LL-Null:                      -1.0639e+05
                                      LLR p-value:                0.1899
=====
              coef      std err          z      P>|z|      [0.025      0.975]
-----
intercept    -1.9888      0.008   -246.669      0.000     -2.005     -1.973
ab_page      -0.0150      0.011    -1.311      0.190     -0.037      0.007
=====
"""
```

- We can see from the above summary that the intercept = -1.9888 and the slope = -0.0150. The p-value of the ab_page is larger than the critical value we would reject the null, in this case, we can say that the specific page does not affect the change in the dependent variable (the coefficient for the specific country is not significant)
- p-value for ab_page = 0.190, in part || the null and the alternative associated with the same conversion rate regardless of the type of page, in the regression model we see the conversion based on which page a customer receives

- Other factors can help in predicting what may affect our response variable, and in trying to find the best model to predict it. No disadvantages came from the additional terms, they provide more evidence relating to our hypothesis

In [207]: *# read in the countries.csv dataset*

```
df3 = pd.read_csv('countries.csv')
df3.head()
```

```
Out[207]:
```

	user_id	country
0	834778	UK
1	928468	US
2	822059	UK
3	711597	UK
4	710616	UK

In [208]: *# merge together the two datasets on the appropriate rows*

```
df4 = df3.set_index('user_id').join(df2.set_index('user_id'), how='inner')
df4.head()
```

```
Out[208]:
```

	country	timestamp	group	landing_page	\
user_id					
834778	UK	2017-01-14 23:08:43.304998	control	old_page	
928468	US	2017-01-23 14:44:16.387854	treatment	new_page	
822059	UK	2017-01-16 14:04:14.719771	treatment	new_page	
711597	UK	2017-01-22 03:14:24.763511	control	old_page	
710616	UK	2017-01-16 13:14:44.000513	treatment	new_page	

	converted	intercept	ab_page
user_id			
834778	0	1	0
928468	0	1	1
822059	1	1	1
711597	0	1	0
710616	0	1	1

In [209]: df4.country.unique()

```
Out[209]: array(['UK', 'US', 'CA'], dtype=object)
```

In [210]: *# create dummy variables on the country columns*

```
df4[['UK', 'US', 'CA']] = pd.get_dummies(df4['country'])
df4.head()
```

```
Out[210]:
```

	country	timestamp	group	landing_page	\
user_id					
834778	UK	2017-01-14 23:08:43.304998	control	old_page	
928468	US	2017-01-23 14:44:16.387854	treatment	new_page	
822059	UK	2017-01-16 14:04:14.719771	treatment	new_page	
711597	UK	2017-01-22 03:14:24.763511	control	old_page	

710616	UK	2017-01-16 13:14:44.000513	treatment	new_page
--------	----	----------------------------	-----------	----------

	converted	intercept	ab_page	UK	US	CA
user_id						
834778	0	1	0	0	1	0
928468	0	1	1	0	0	1
822059	1	1	1	0	1	0
711597	0	1	0	0	1	0
710616	0	1	1	0	1	0

In [211]: *# Delete one column of the three dummy variables*

```
df4 = df4.drop('CA', axis = 1)
df4.head()
```

Out[211]:

	country	timestamp	group	landing_page \
--	---------	-----------	-------	----------------

user_id				
834778	UK	2017-01-14 23:08:43.304998	control	old_page
928468	US	2017-01-23 14:44:16.387854	treatment	new_page
822059	UK	2017-01-16 14:04:14.719771	treatment	new_page
711597	UK	2017-01-22 03:14:24.763511	control	old_page
710616	UK	2017-01-16 13:14:44.000513	treatment	new_page

	converted	intercept	ab_page	UK	US
user_id					
834778	0	1	0	0	1
928468	0	1	1	0	0
822059	1	1	1	0	1
711597	0	1	0	0	1
710616	0	1	1	0	1

In [212]: *# fit the regression model*

```
model2 = sm.Logit(df4['converted'], df4[['intercept', 'UK', 'US', 'ab_page']])
result2 = model2.fit()
result2.summary()
```

Optimization terminated successfully.

Current function value: 0.366113

Iterations 6

Out[212]: <class 'statsmodels.iolib.summary.Summary'>

```
"""
                                Logit Regression Results
=====
Dep. Variable:                  converted    No. Observations:                  290584
Model:                            Logit      Df Residuals:                      290580
Method:                           MLE        Df Model:                          3
Date:                Tue, 22 Jan 2019    Pseudo R-squ.:                    2.323e-05
Time:                      15:22:28      Log-Likelihood:                   -1.0639e+05
```

```

converged: True LL-Null: -1.0639e+05
LLR p-value: 0.1760
=====
              coef      std err          z      P>|z|      [0.025      0.975]
-----
intercept    -1.9893      0.009   -223.763      0.000     -2.007     -1.972
UK           -0.0408      0.027    -1.516      0.130     -0.093      0.012
US            0.0099      0.013      0.743      0.457     -0.016      0.036
ab_page      -0.0149      0.011    -1.307      0.191     -0.037      0.007
=====
"""

```

- We can see that for the US we predict an increase in conversion by 0.0099 holding other variables constant. For every 100% increase in the UK, the predicted decrease in conversion is -0.0408.
- The only p-value that is statistically significant is the intercept and in this case, we can say that a specific country does not affect the change in the conversion**

Now we will look at an interaction between page and country to see if there significant effects on conversion

```

In [213]: # Create `ab_test_page` column, which is 1 if `new_page` and 0 if `old_page`
df4['ab_test_page'] = (df4['landing_page'] == "new_page").astype(int)
df4.head()

```

```

Out[213]:
   country  timestamp  group landing_page \
user_id
834778    UK  2017-01-14 23:08:43.304998  control  old_page
928468    US  2017-01-23 14:44:16.387854  treatment  new_page
822059    UK  2017-01-16 14:04:14.719771  treatment  new_page
711597    UK  2017-01-22 03:14:24.763511  control  old_page
710616    UK  2017-01-16 13:14:44.000513  treatment  new_page

   converted  intercept  ab_page  UK  US  ab_test_page
user_id
834778      0          1        0  0  1              0
928468      0          1        1  0  0              1
822059      1          1        1  0  1              1
711597      0          1        0  0  1              0
710616      0          1        1  0  1              1

```

```

In [214]: # Delete one column of the three dummy variables
#df4 = df4.drop('CA', axis = 1)
df4.head()

```

```

Out[214]:
   country  timestamp  group landing_page \
user_id
834778    UK  2017-01-14 23:08:43.304998  control  old_page
928468    US  2017-01-23 14:44:16.387854  treatment  new_page

```

822059	UK	2017-01-16 14:04:14.719771	treatment	new_page
711597	UK	2017-01-22 03:14:24.763511	control	old_page
710616	UK	2017-01-16 13:14:44.000513	treatment	new_page

	converted	intercept	ab_page	UK	US	ab_test_page
user_id						
834778	0	1	0	0	1	0
928468	0	1	1	0	0	1
822059	1	1	1	0	1	1
711597	0	1	0	0	1	0
710616	0	1	1	0	1	1

```
In [215]: # fit the regression model
model3 = sm.Logit(df4['ab_test_page'], df4[['intercept', 'UK', 'US']])
result3 = model3.fit()
result3.summary()
```

Optimization terminated successfully.
Current function value: 0.693144
Iterations 3

```
Out[215]: <class 'statsmodels.iolib.summary.Summary'>
"""
```

```

                                Logit Regression Results
=====
Dep. Variable:                ab_test_page    No. Observations:                290584
Model:                            Logit      Df Residuals:                290581
Method:                            MLE       Df Model:                      2
Date:                Tue, 22 Jan 2019      Pseudo R-squ.:                4.442e-06
Time:                        15:22:29      Log-Likelihood:               -2.0142e+05
converged:                        True      LL-Null:                      -2.0142e+05
                                      LLR p-value:                0.4088
=====
               coef      std err          z      P>|z|      [0.025      0.975]
-----
intercept      0.0018      0.004      0.414      0.679      -0.007      0.011
UK              0.0124      0.017      0.720      0.472      -0.021      0.046
US             -0.0088      0.009     -1.023      0.306      -0.026      0.008
=====
"""
```

- We can see that for the UK we predict an increase in conversion by 0.0124 holding other variables constant. For every 100% increase in the US, the predicted decrease in conversion is -0.0088
- p-values are not statistically significant, and the coefficient for the specific country is not significant

0.2.4 Conclusions

- Considering all these factors along with the timestamp for this experiment, we can conclude that the duration of this experiment is not enough to decide whether the new page is better or not. Country and Page factors do not affect the change in the conversion rate and the coefficients are not significant and so, we fail to reject the null hypothesis. The company needs to run the experiment longer to make their decision.